

References

CS231n: Convolutional Neural Networks for Visual Recognition



Fei-Fei Li & Andrej Karpathy & Justin Johnson

Fast R-CNN

Rich feature hierarchies for accurate object detection and semantic segmentation

Ross Girshick
Microsoft Research
rbg@microsoft.com

Ross Girshick Jeff Donahue Trevor Darrell Jitendra Malik
UC Berkeley
{rbg,jdonahue,trevor,malik}@eecs.berkeley.edu

**Faster R-CNN: Towards Real-Time Object
Detection with Region Proposal Networks**

Shaoqing Ren, Kaiming He, Ross Girshick, and Jian Sun

**Spatial Pyramid Pooling in Deep Convolutional
Networks for Visual Recognition**

Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun

Selective Search for Object Recognition

J.R.R. Uijlings^{*1,2}, K.E.A. van de Sande^{†2}, T. Gevers², and A.W.M. Smeulders²

¹University of Trento, Italy



²University of Amsterdam, the Netherlands

Technical Report 2012, submitted to IJCV

Fully Convolutional Networks for Semantic Segmentation

Jonathan Long^{*} Evan Shelhamer^{*} Trevor Darrell
UC Berkeley
{jonlong,shelhamer,trevor}@cs.berkeley.edu


“A dramatic moment in the meteoric rise of deep learning came when a convolutional network won this challenge for the first time and by a wide margin, bringing down the state-of-the-art top-5 error rate from 26.1% to 15.3% (Krizhevsky *et al.*, 2012), meaning that the convolutional network produces a ranked list of possible categories for each image and the correct category appeared in the first five entries of this list for all but 15.3% of the test examples. Since then, these competitions are consistently won by deep convolutional nets, and as of this writing, advances in deep learning have brought the latest top-5 error rate in this contest down to 3.6%” – Ref: Deep Learning Book by Y Bengio et al



What is a convolutional neural network?

Convolutional networks are simply neural networks that use convolution in place of general matrix multiplication in at least one of their layers.

- Convolution is a mathematical operation having a linear form



Types of inputs

- Inputs have a structure
 - Color images are three dimensional and so have a volume
 - Time domain speech signals are 1-d while the frequency domain representations (e.g. MFCC vectors) take a 2d form. They can also be looked at as a time sequence.
 - Medical images (such as CT/MR/etc) are multidimensional
 - Videos have the additional temporal dimension compared to stationary images
 - Speech signals can be modelled as 2 dimensional
 - Variable length sequences and time series data are again multidimensional
- Hence it makes sense to **model them as tensors** instead of vectors.
- The classifier then needs to accept a tensor as input and perform the necessary machine learning task. In the case of an image, this tensor represents a volume.

CNNs are everywhere

- Image retrieval
- Detection
- Self driving cars
- Semantic segmentation
- Face recognition (FB tagging)
- Pose estimation
- Detect diseases
- Speech Recognition
- Text processing
- Analysing satellite data

CNNs for applications that involve images

- Why CNNs are more suitable to process images?
- Pixels in an image correlate to each other. However, nearby pixels correlate stronger and distant pixels don't influence much
 - Local features are important: Local Receptive Fields
- Affine transformations: The class of an image doesn't change with translation. We can build a feature detector that can look for a particular feature (e.g. an edge) anywhere in the image plane by moving across. A convolutional layer may have several such filters constituting the depth dimension of the layer.

Fully connected layers

- Fully connected layers (such as the hidden layers of a traditional neural network) are agnostic to the structure of the input
 - They take inputs as vectors and generate an output vector
 - There is no requirement to share parameters unless forced upon in specific architectures. This blows up the number of parameters as the input and/or output dimensions increase.
- Suppose we are to perform classification on an image of 100x100x3 dimensions.
- If we implement using a feed forward neural network that has an input, hidden and an output layer, where: hidden units (nh) = 1000, output classes = 10 :
 - Input layer = 10k pixels * 3 = 30k, weight matrix for hidden to input layer = 1k * 30k = 30 M and output layer matrix size = 10 * 1000 = 10k
- We may handle this is by extracting the features using pre processing and presenting a lower dimensional input to the Neural Network. But this requires expert engineered features and hence domain knowledge

Convolution

Convolution in 1 Dimension:

$$y[n] = \sum_{k=-\infty}^{k=\infty} x[k]h[n-k]$$

Convolution in 2 Dimensions:

$$y[n_1, n_2] = \sum_{k_1=-\infty}^{k_1=\infty} \sum_{k_2=-\infty}^{k_2=\infty} x[k_1, k_2]h[(n_1 - k_1), (n_2 - k_2)]$$

Input

| | | | |
|-----|-----|-----|-----|
| a | b | c | d |
| e | f | g | h |
| i | j | k | l |

Kernel

| | |
|-----|-----|
| w | x |
| y | z |

Output

$$\begin{array}{rcl} aw & + & bx \\ ey & + & fz \end{array} +$$

$$\begin{array}{rcl} bw & + & cx \\ fy & + & gz \end{array} +$$

$$\begin{array}{rcl} cw & + & dx \\ gy & + & hz \end{array} +$$

$$\begin{array}{rcl} ew & + & fx \\ iy & + & jz \end{array} +$$

$$\begin{array}{rcl} fw & + & gx \\ jy & + & kz \end{array} +$$

$$\begin{array}{rcl} gw & + & hx \\ ky & + & lz \end{array} +$$


CNNs

Types of layers in a CNN:

- Convolution Layer
- Pooling Layer
- Fully Connected Layer

Case Study: LeNet-5

[LeCun et al., 1998]




Conv filters were 5x5, applied at stride 1

Subsampling (Pooling) layers were 2x2 applied at stride 2
i.e. architecture is [CONV-POOL-CONV-POOL-CONV-FC]

Convolution Layer


- A layer in a regular neural network take vector as input and output a vector.

Regular neural network (fully connected):



- A convolution layer takes a tensor (3d volume for RGB images) as input and generates a tensor as output


Convolutional neural network:



Each layer takes a 3d volume, produces 3d volume with some smooth function that may or may not have parameters.

Convolutional Neural Networks: Layers



- **INPUT** [32x32x3] will hold the raw pixel values of the image, in this case an image of width 32, height 32, and with three color channels R,G,B.
- **CONV** layer will compute the output of neurons that are connected to local regions in the input, each computing a dot product between their weights and a small region they are connected to in the input volume. This may result in volume such as [32x32x12] if we decided to use 12 filters.
- **RELU** layer will apply an elementwise activation function, such as the $\max(0,x)$ thresholding at zero. This leaves the size of the volume unchanged ([32x32x12]).
- **POOL** layer will perform a downsampling operation along the spatial dimensions (width, height), resulting in volume such as [16x16x12].
- **FC** (i.e. fully-connected) layer will compute the class scores, resulting in volume of size [1x1x10], where each of the 10 numbers correspond to a class score, such as among the 10 categories of CIFAR-10. As with ordinary Neural Networks and as the name implies, each neuron in this layer will be connected to all the numbers in the previous volume.



Local Receptive Fields

- Filter (Kernel) is applied on the input image like a moving window along width and height
- The depth of a filter matches that of the input.
- For each position of the filter, the dot product of filter and the input are computed (Activation)
- The 2d arrangement of these activations is called an activation map.
- The number of such filters constitute the depth of the convolution layer

Dealing with Images: Local Connectivity




Same neuron. Just more focused (narrow “receptive field”).

The parameters on a each filter are spatially “shared”
(if a feature is useful in one place, it’s useful elsewhere)

Convolution Operation between filter and image

- The convolution layer computes dot products between the filter and a piece of image as it slides along the image
- The step size of slide is called stride
- Without any padding, the convolution process decreases the spatial dimensions of the output




Activation Maps

- Example:
 - Consider an image $32 \times 32 \times 3$ and a $5 \times 5 \times 3$ filter.
 - The convolution happens between a $5 \times 5 \times 3$ chunk of the image with the filter: $w^T x + b$
 - In this example we get 75 dimensional vector and a bias term
 - In this example, with a stride of 1, we get $28 \times 28 \times 1$ activation for 1 filter without padding
 - If we have 6 filters, we would get $28 \times 28 \times 6$ without padding
- In the above example we have an activation map of 28×28 per filter.
- Activation maps are feature inputs to the subsequent layer of the network
- Without any padding, the 2D surface area of the activation map is smaller than the input surface area for a stride of ≥ 1

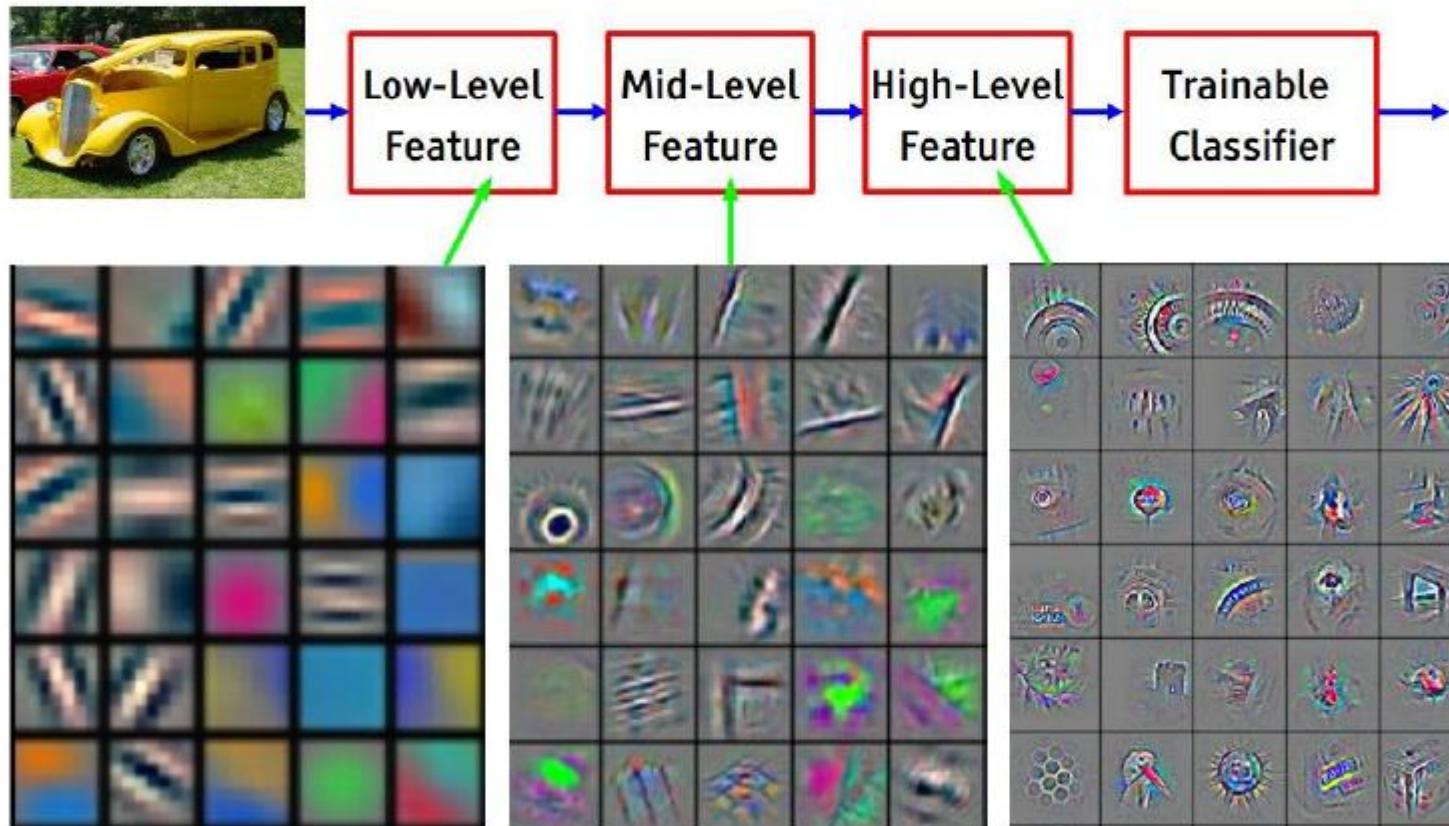
Stacking Convolution Layers

Preview: ConvNet is a sequence of Convolutional Layers, interspersed with activation functions



Feature Representation as a hierarchy

*[From recent Yann
LeCun slides]*



Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

Padding

- The spatial (x, y) extent of the output produced by the convolutional layer is less than the respective dimensions of the input (except for the special case of 1 x 1 filter with a stride 1).
- As we add more layers and use larger strides, the output surface dimensions keep reducing and this may impact the accuracy.
- Often, we may want to preserve the spatial extent during the initial layers and downsample them at a later time.
- Padding the input with suitable values (padding with zero is common) helps to preserve the spatial size

Zero Padding the border

| | | | | | | | | |
|---|---|---|---|---|---|--|--|--|
| 0 | 0 | 0 | 0 | 0 | 0 | | | |
| 0 | | | | | | | | |
| 0 | | | | | | | | |
| 0 | | | | | | | | |
| 0 | | | | | | | | |
| | | | | | | | | |
| | | | | | | | | |
| | | | | | | | | |
| | | | | | | | | |

e.g. input 7x7

3x3 filter, applied with **stride 1**

pad with 1 pixel border => what is the output?

7x7 output!

in general, common to see CONV layers with stride 1, filters of size $F \times F$, and zero-padding with $(F-1)/2$. (will preserve size spatially)

e.g. $F = 3 \Rightarrow$ zero pad with 1

$F = 5 \Rightarrow$ zero pad with 2

$F = 7 \Rightarrow$ zero pad with 3

Hyperparameters of the convolution layer

- Filter Size

Common settings:

K = (powers of 2, e.g. 32, 64, 128, 512)

- $F = 3, S = 1, P = 1$
- $F = 5, S = 1, P = 2$
- $F = 5, S = 2, P = ?$ (whatever fits)
- $F = 1, S = 1, P = 0$

- # Filters

Summary. To summarize, the Conv Layer:

- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires four hyperparameters:
 - Number of filters K ,
 - their spatial extent F ,
 - the stride S ,
 - the amount of zero padding P .
- Produces a volume of size $W_2 \times H_2 \times D_2$ where:
 - $W_2 = (W_1 - F + 2P)/S + 1$
 - $H_2 = (H_1 - F + 2P)/S + 1$ (i.e. width and height are computed equally by symmetry)
 - $D_2 = K$
- With parameter sharing, it introduces $F \cdot F \cdot D_1$ weights per filter, for a total of $(F \cdot F \cdot D_1) \cdot K$ weights and K biases.
- In the output volume, the d -th depth slice (of size $W_2 \times H_2$) is the result of performing a valid convolution of the d -th filter over the input volume with a stride of S , and then offset by d -th bias.

- Stride


- Padding

Pooling Layer

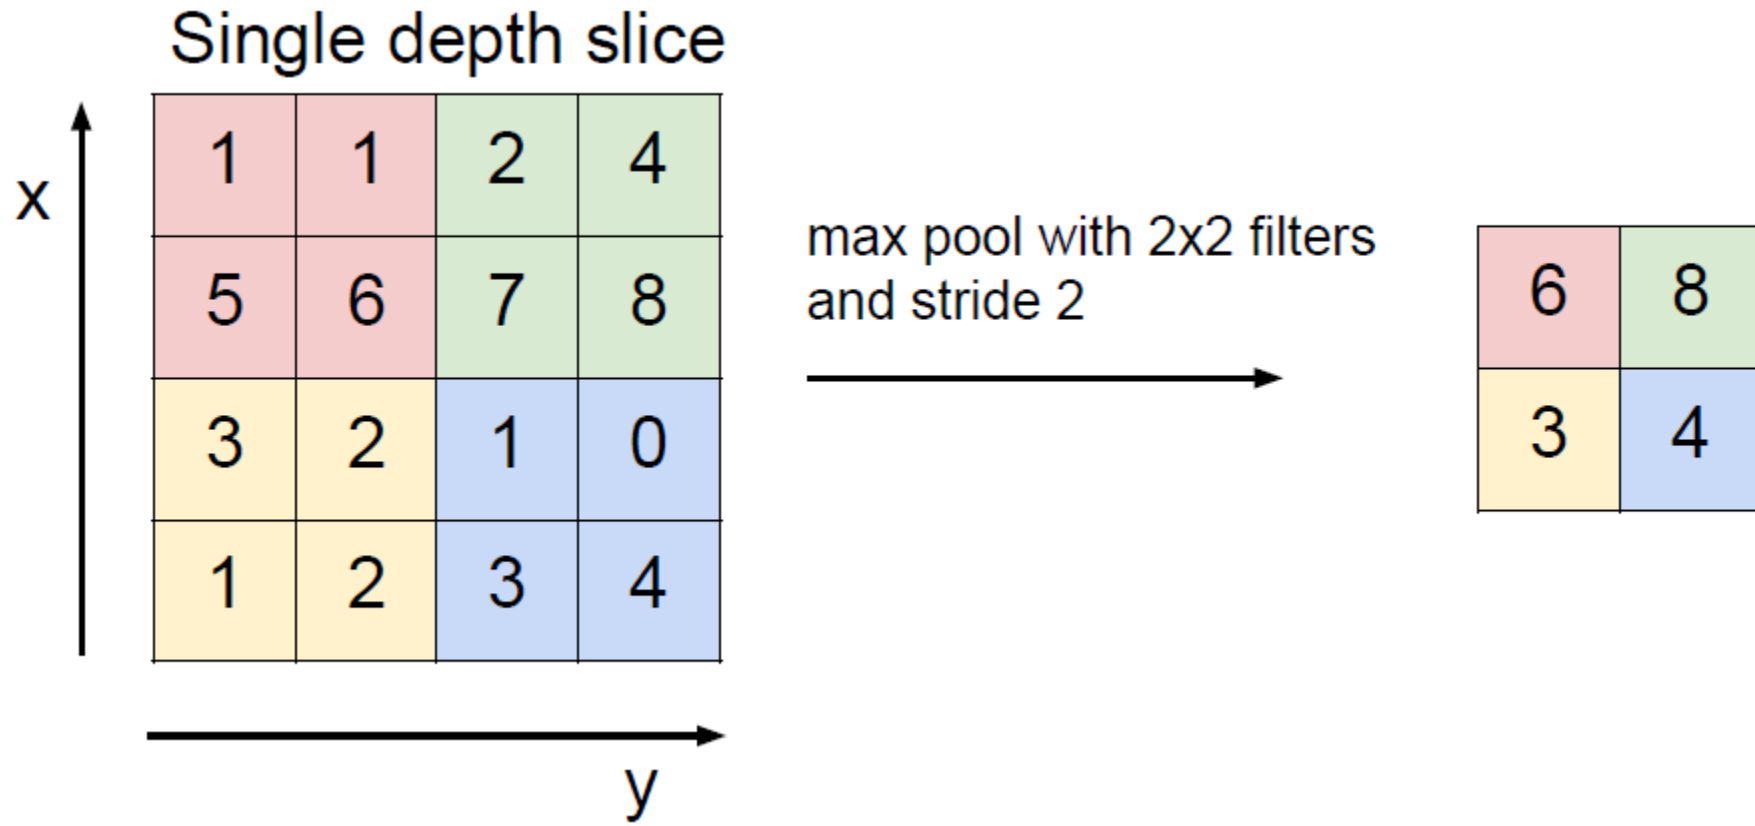
- Pooling is a downsampling operation
- The rationale is that the “meaning” embedded in a piece of image can be captured using a small subset of “important” pixels
- Max pooling and average pooling are the two most common operations
- Pooling layer doesn’t have any trainable parameters

Pooling layer

- makes the representations smaller and more manageable
- operates over each activation map independently:




Max Pooling Illustration



Popular Network Architectures

Current trend: Deeper Models

- CNNs consistently outperform other approaches for the core tasks of CV
- Deeper models work better
- Increasing the number of parameters in layers of CNN without increasing their depth is not effective at increasing test set performance.
- Shallow models overfit at around 20 million parameters while deep ones can benefit from having over 60 million.
- Key insight: Model performs better when it is architected to reflect composition of simpler functions than a single complex function. This may also be explained off viewing the computation as a chain of dependencies



Case Study: AlexNet

[Krizhevsky et al. 2012]

Full (simplified) AlexNet architecture:

[224x224x3] INPUT

[55x55x96] **CONV1**: 96 11x11 filters at stride 4, pad 0

[27x27x96] **MAX POOL1**: 3x3 filters at stride 2

[27x27x96] **NORM1**: Normalization layer

[27x27x256] **CONV2**: 256 5x5 filters at stride 1, pad 2

[13x13x256] **MAX POOL2**: 3x3 filters at stride 2

[13x13x256] **NORM2**: Normalization layer

[13x13x384] **CONV3**: 384 3x3 filters at stride 1, pad 1

[13x13x384] **CONV4**: 384 3x3 filters at stride 1, pad 1


[13x13x256] **CONV5**: 256 3x3 filters at stride 1, pad 1

[6x6x256] **MAX POOL3**: 3x3 filters at stride 2

[4096] **FC6**: 4096 neurons

[4096] **FC7**: 4096 neurons

[1000] **FC8**: 1000 neurons (class scores)



Details/Retrospectives:

- first use of ReLU
- used Norm layers (not common anymore)
- heavy data augmentation
- dropout 0.5
- batch size 128
- SGD Momentum 0.9
- Learning rate 1e-2, reduced by 10 manually when val accuracy plateaus
- L2 weight decay 5e-4
- 7 CNN ensemble: 18.2% -> 15.4%

VGG Net

INPUT: [224x224x3] memory: $224*224*3=150K$ params: 0 (not counting biases)

CONV3-64: [224x224x64] memory: $224*224*64=3.2M$ params: $(3*3*3)*64 = 1,728$

CONV3-64: [224x224x64] memory: $224*224*64=3.2M$ params: $(3*3*64)*64 = 36,864$

POOL2: [112x112x64] memory: $112*112*64=800K$ params: 0

CONV3-128: [112x112x128] memory: $112*112*128=1.6M$ params: $(3*3*64)*128 = 73,728$

CONV3-128: [112x112x128] memory: $112*112*128=1.6M$ params: $(3*3*128)*128 = 147,456$

POOL2: [56x56x128] memory: $56*56*128=400K$ params: 0

CONV3-256: [56x56x256] memory: $56*56*256=800K$ params: $(3*3*128)*256 = 294,912$

CONV3-256: [56x56x256] memory: $56*56*256=800K$ params: $(3*3*256)*256 = 589,824$

CONV3-256: [56x56x256] memory: $56*56*256=800K$ params: $(3*3*256)*256 = 589,824$

POOL2: [28x28x256] memory: $28*28*256=200K$ params: 0

CONV3-512: [28x28x512] memory: $28*28*512=400K$ params: $(3*3*256)*512 = 1,179,648$

CONV3-512: [28x28x512] memory: $28*28*512=400K$ params: $(3*3*512)*512 = 2,359,296$

CONV3-512: [28x28x512] memory: $28*28*512=400K$ params: $(3*3*512)*512 = 2,359,296$

POOL2: [14x14x512] memory: $14*14*512=100K$ params: 0

CONV3-512: [14x14x512] memory: $14*14*512=100K$ params: $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] memory: $14*14*512=100K$ params: $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] memory: $14*14*512=100K$ params: $(3*3*512)*512 = 2,359,296$

POOL2: [7x7x512] memory: $7*7*512=25K$ params: 0

FC: [1x1x4096] memory: 4096 params: $7*7*512*4096 = 102,760,448$

FC: [1x1x4096] memory: 4096 params: $4096*4096 = 16,777,216$

FC: [1x1x1000] memory: 1000 params: $4096*1000 = 4,096,000$

TOTAL memory: $24M * 4 \text{ bytes} \approx 93MB$ / image (only forward! $\sim *2$ for bwd)

TOTAL params: 138M parameters

| ConvNet Configuration | | |
|---------------------------|------------------|------------------|
| B | C | D |
| 13 weight layers | 16 weight layers | 16 weight layers |
| put (224 × 224 RGB image) | | |
| conv3-64 | conv3-64 | conv3-64 |
| conv3-64 | conv3-64 | conv3-64 |
| maxpool | | |
| conv3-128 | conv3-128 | conv3-128 |
| conv3-128 | conv3-128 | conv3-128 |
| maxpool | | |
| conv3-256 | conv3-256 | conv3-256 |
| conv3-256 | conv3-256 | conv3-256 |
| maxpool | | |
| conv3-512 | conv3-512 | conv3-512 |
| conv3-512 | conv3-512 | conv3-512 |
| maxpool | | |
| conv3-512 | conv3-512 | conv3-512 |
| conv3-512 | conv3-512 | conv3-512 |
| maxpool | | |
| conv3-512 | conv3-512 | conv3-512 |
| conv3-512 | conv3-512 | conv3-512 |
| maxpool | | |
| FC-4096 | | |
| FC-4096 | | |
| FC-1000 | | |
| soft-max | | |

VGG net

INPUT: [224x224x3] memory: $224*224*3=150\text{K}$ params: 0 (not counting biases)

CONV3-64: [224x224x64] memory: $224*224*64=3.2\text{M}$ params: $(3*3*3)*64 = 1,728$

CONV3-64: [224x224x64] memory: $224*224*64=3.2\text{M}$ params: $(3*3*64)*64 = 36,864$

POOL2: [112x112x64] memory: $112*112*64=800\text{K}$ params: 0

CONV3-128: [112x112x128] memory: $112*112*128=1.6\text{M}$ params: $(3*3*64)*128 = 73,728$

CONV3-128: [112x112x128] memory: $112*112*128=1.6\text{M}$ params: $(3*3*128)*128 = 147,456$

POOL2: [56x56x128] memory: $56*56*128=400\text{K}$ params: 0

CONV3-256: [56x56x256] memory: $56*56*256=800\text{K}$ params: $(3*3*128)*256 = 294,912$

CONV3-256: [56x56x256] memory: $56*56*256=800\text{K}$ params: $(3*3*256)*256 = 589,824$

CONV3-256: [56x56x256] memory: $56*56*256=800\text{K}$ params: $(3*3*256)*256 = 589,824$

POOL2: [28x28x256] memory: $28*28*256=200\text{K}$ params: 0

CONV3-512: [28x28x512] memory: $28*28*512=400\text{K}$ params: $(3*3*256)*512 = 1,179,648$

CONV3-512: [28x28x512] memory: $28*28*512=400\text{K}$ params: $(3*3*512)*512 = 2,359,296$

CONV3-512: [28x28x512] memory: $28*28*512=400\text{K}$ params: $(3*3*512)*512 = 2,359,296$

POOL2: [14x14x512] memory: $14*14*512=100\text{K}$ params: 0

CONV3-512: [14x14x512] memory: $14*14*512=100\text{K}$ params: $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] memory: $14*14*512=100\text{K}$ params: $(3*3*512)*512 = 2,359,296$

CONV3-512: [14x14x512] memory: $14*14*512=100\text{K}$ params: $(3*3*512)*512 = 2,359,296$

POOL2: [7x7x512] memory: $7*7*512=25\text{K}$ params: 0

FC: [1x1x4096] memory: 4096 params: $7*7*512*4096 = 102,760,448$

FC: [1x1x4096] memory: 4096 params: $4096*4096 = 16,777,216$

FC: [1x1x1000] memory: 1000 params: $4096*1000 = 4,096,000$

Note:


Most memory is in
early CONV

Most params are
in late FC

TOTAL memory: $24\text{M} * 4 \text{ bytes} \approx 93\text{MB}$ / image (only forward! $\sim *2$ for bwd)

TOTAL params: 138M parameters

ResNet



(slide from Kaiming He's recent presentation)

Core Tasks of Computer Vision

| Core CV Task | Task Description | Output | Metrics |
|-----------------------|--|-------------------------------|--|
| Classification | Given an image, assign a label | Class Label | Accuracy |
| Localization | Determine the bounding box containing the object in the given image | Box given by (x1, y1, x2, y2) | Ratio of intersection to the union (Overlap) between the ground truth and bounding box |
| Object Detection | Given an image, detect all the objects and their locations in the image | For each object: (Label, Box) | Mean Avg Best Overlap (MABO,) mean Average Precision (mAP) |
| Semantic Segmentation | Given an image, assign each pixel to a class label, so that we can look at the image as a set of labelled segments | A set of image segments | Classification metrics, Intersection by Union overlap |
| Instance Segmentation | Same as semantic segmentation, but each instance of a segment class is determined uniquely | A set of image segments | |

Object Localization

- Given an image containing an object of interest, determine the bounding box for the object
- Classify the object

**Classification
+ Localization**



Classification + Localization: Task

Classification: C classes

Input: Image

Output: Class label

Evaluation metric: Accuracy



CAT

Localization:

Input: Image

Output: Box in the image (x, y, w, h)

Evaluation metric: Intersection over Union



(x, y, w, h)

Classification + Localization: Do both

Classification + Localization: ImageNet

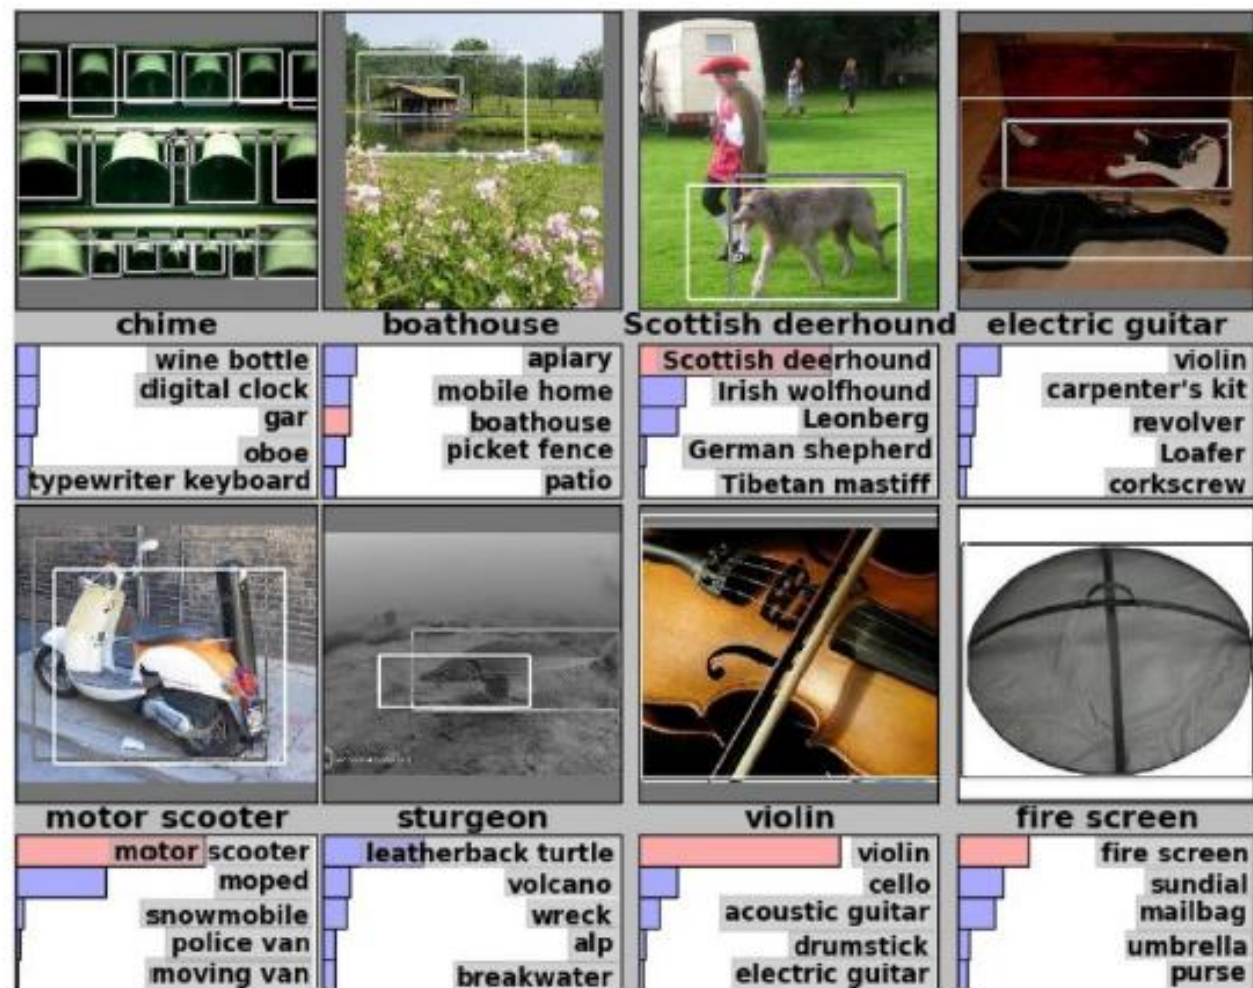
1000 classes (same as classification)

Each image has 1 class, at least one bounding box

~800 training images per class

Algorithm produces 5 (class, box) guesses

Example is correct if at least one guess has correct class AND bounding box at least 0.5 intersection over union (IoU)



Idea #1: Localization as Regression

Input: image



Only one object,
simpler than detection

Neural Net
→

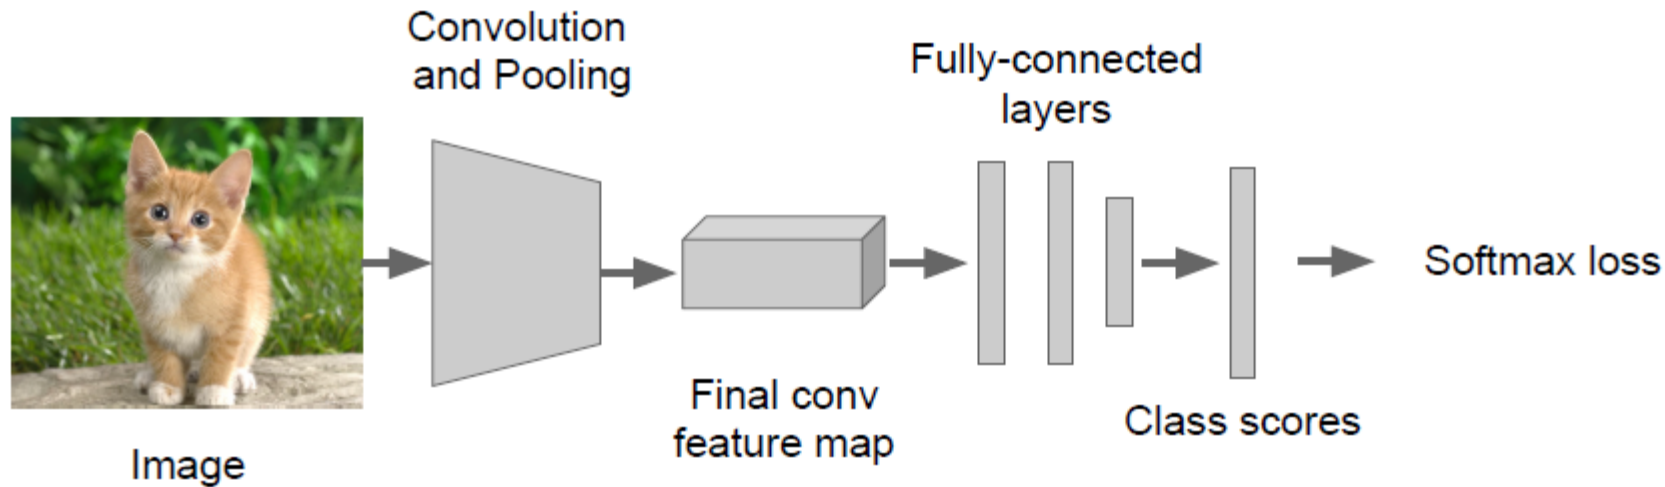
Output:
Box coordinates
(4 numbers)

Correct output:
box coordinates
(4 numbers)

Loss:
L2 distance

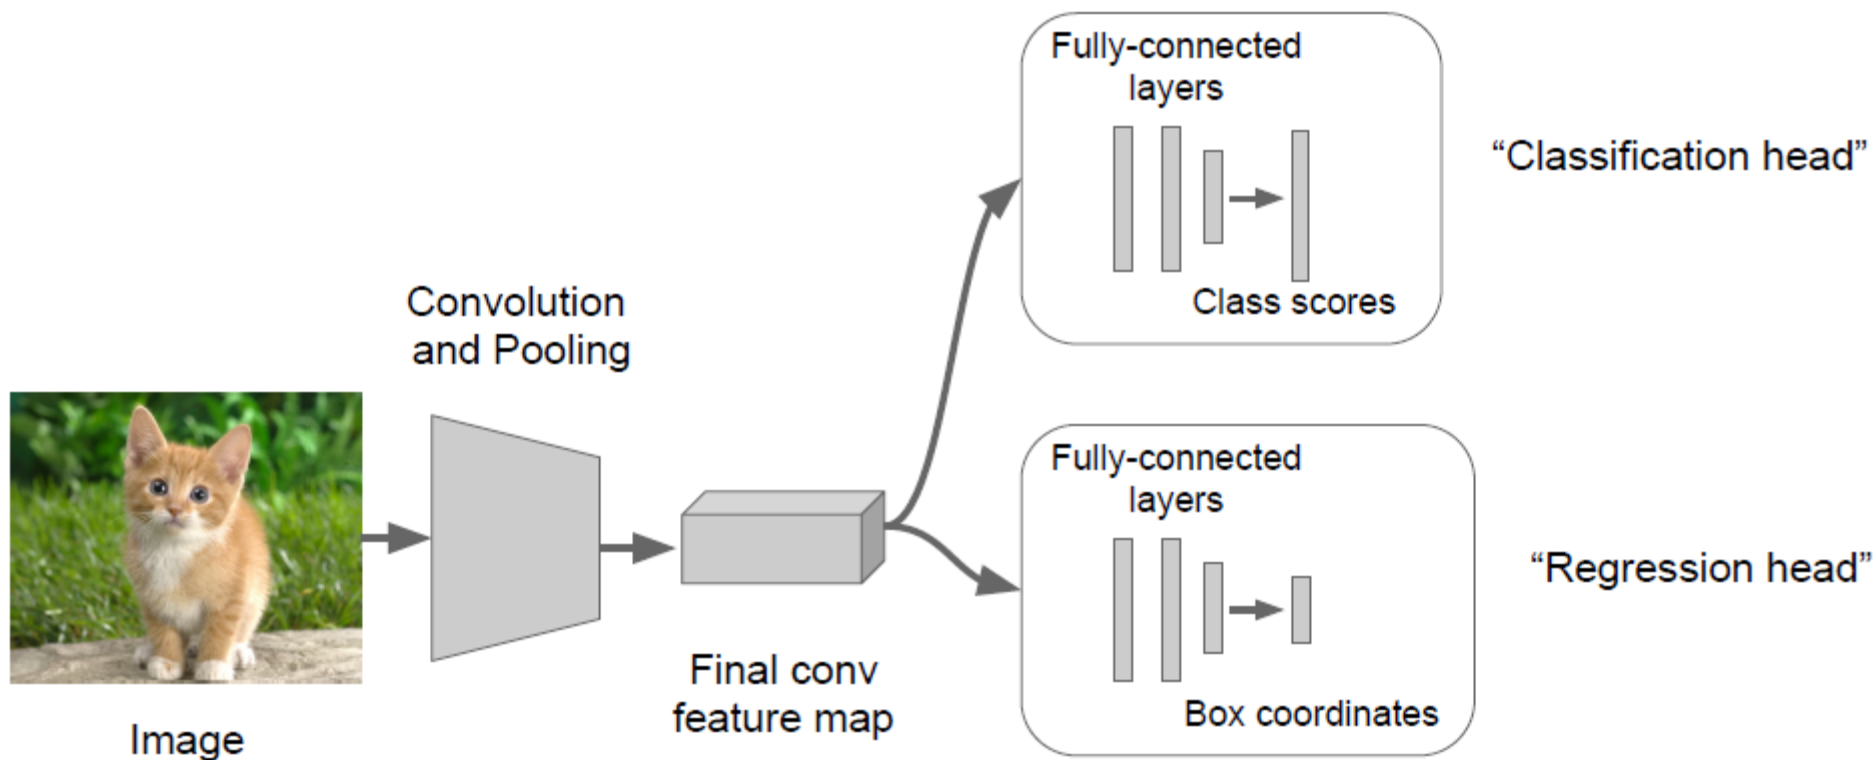
Simple Recipe for Classification + Localization

Step 1: Train (or download) a classification model (AlexNet, VGG, GoogLeNet)



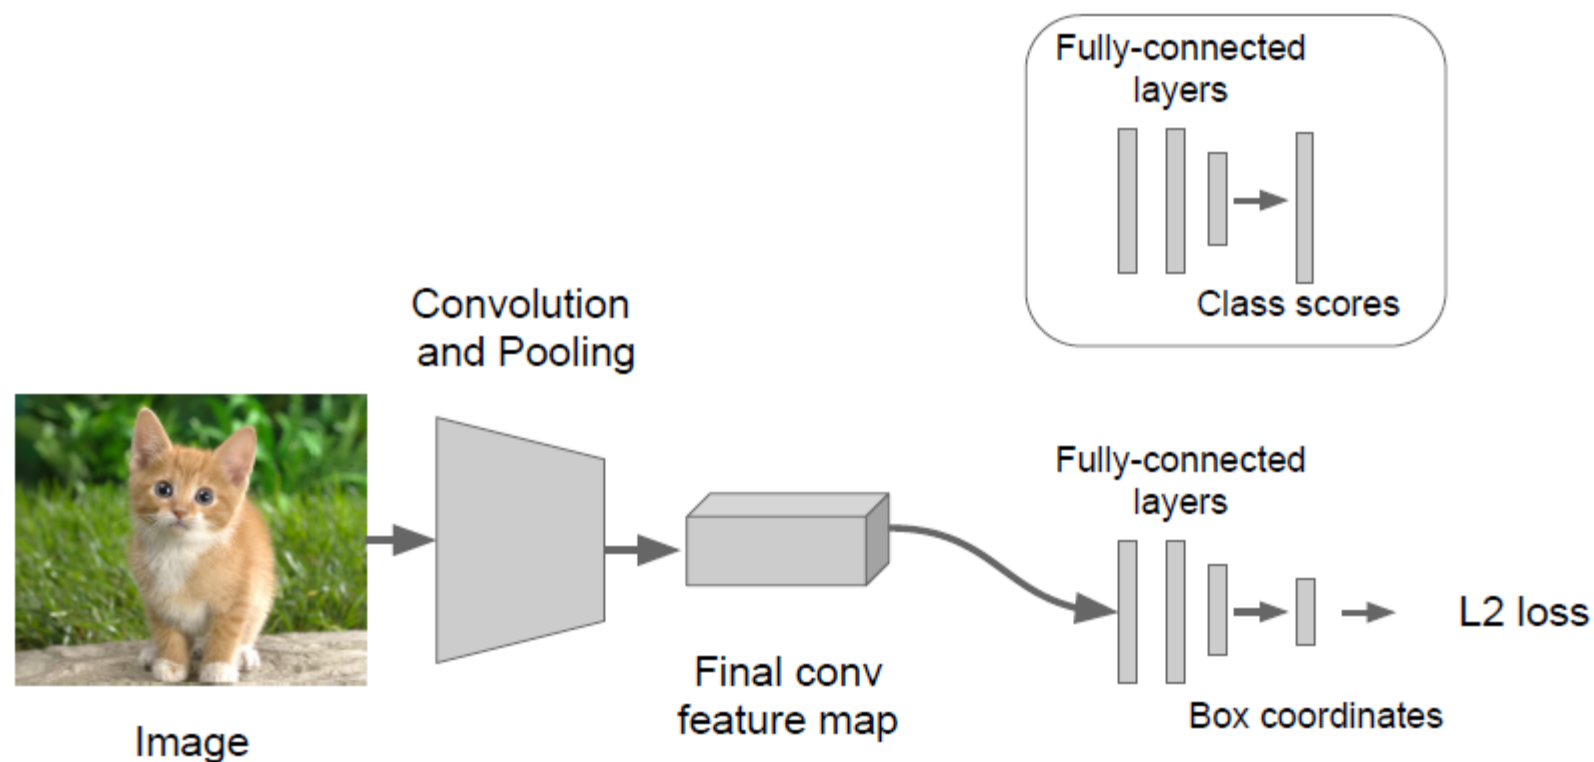
Simple Recipe for Classification + Localization

Step 2: Attach new fully-connected “regression head” to the network



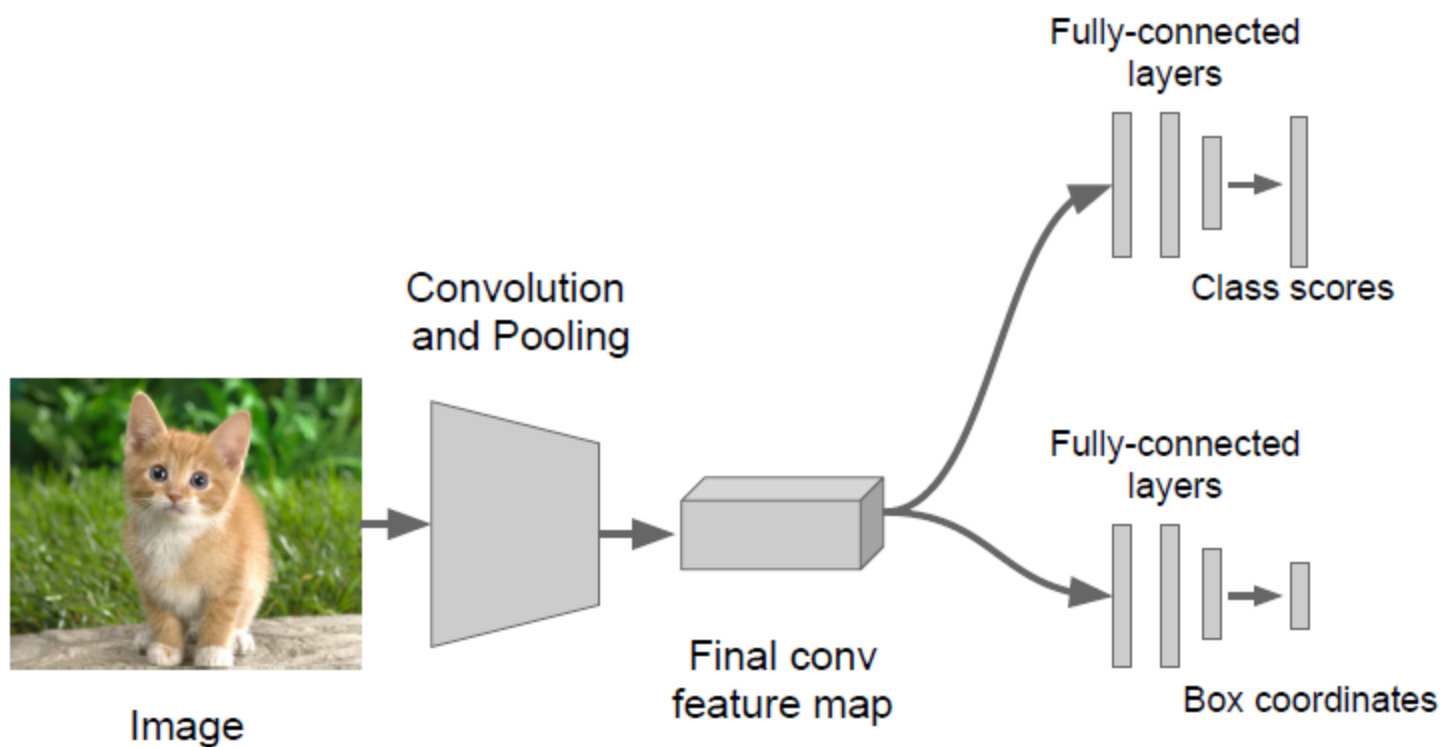
Simple Recipe for Classification + Localization

Step 3: Train the regression head only with SGD and L2 loss



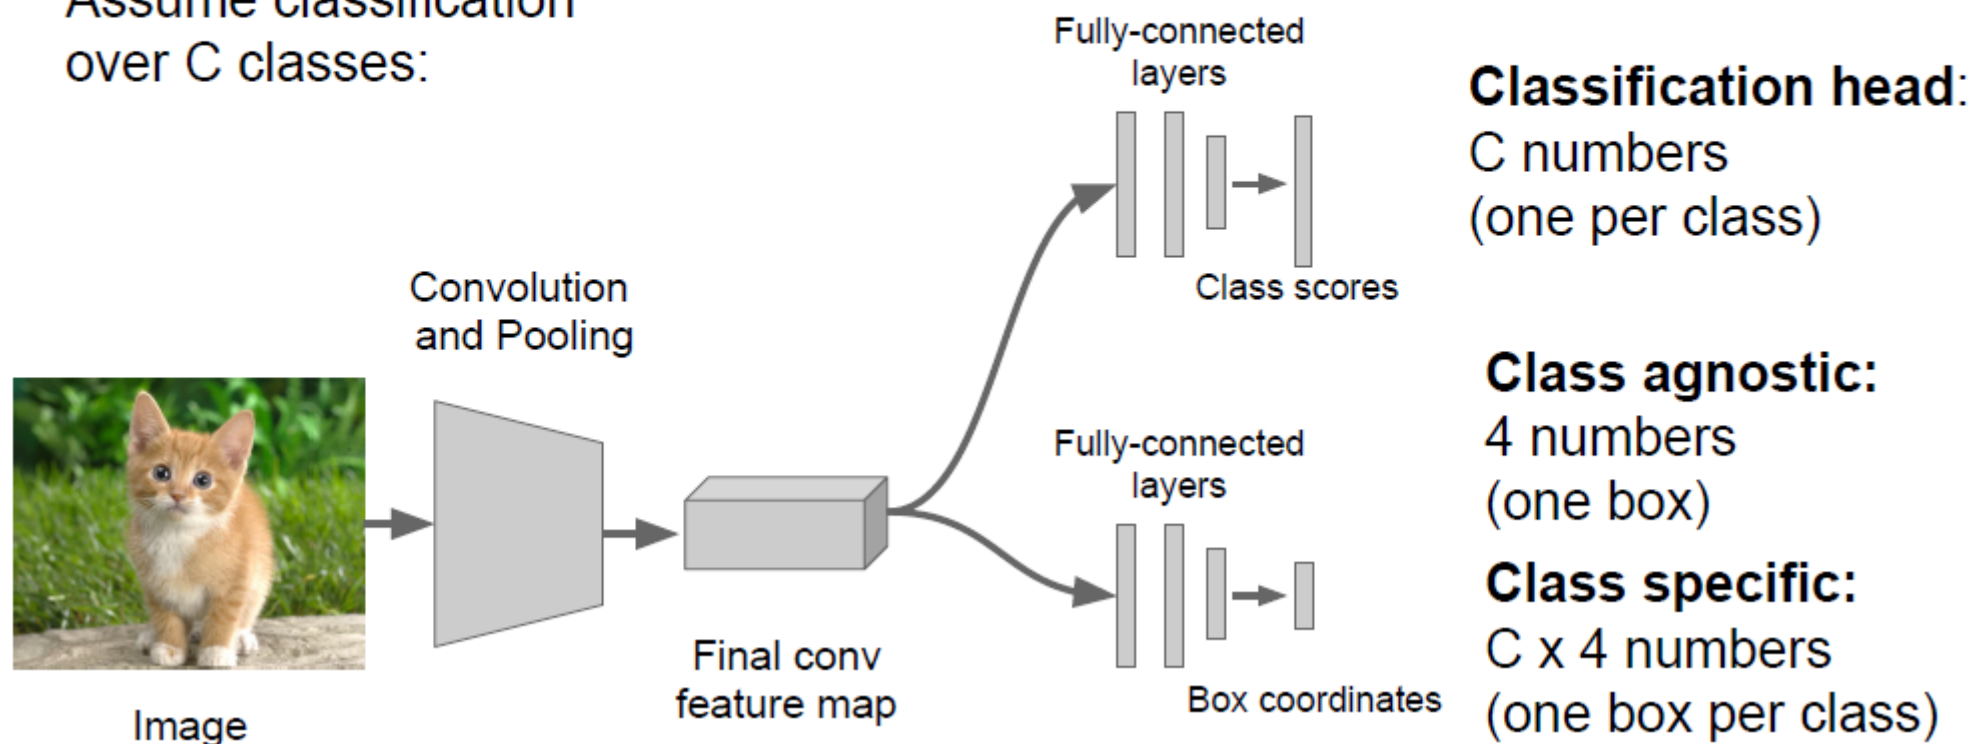
Simple Recipe for Classification + Localization

Step 4: At test time use both heads

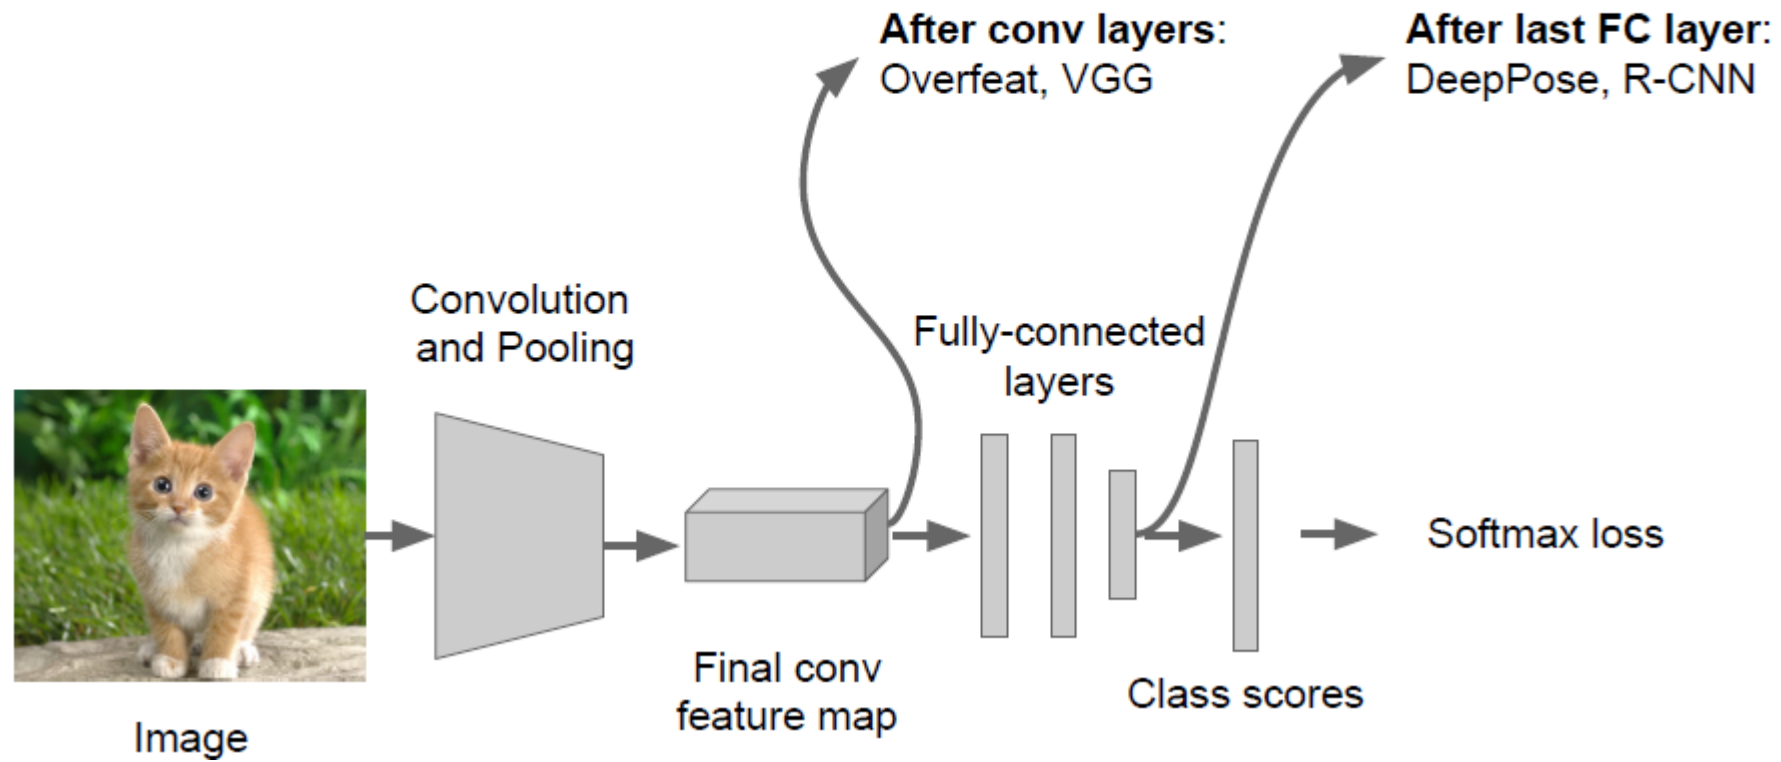


Per-class vs class agnostic regression

Assume classification
over C classes:



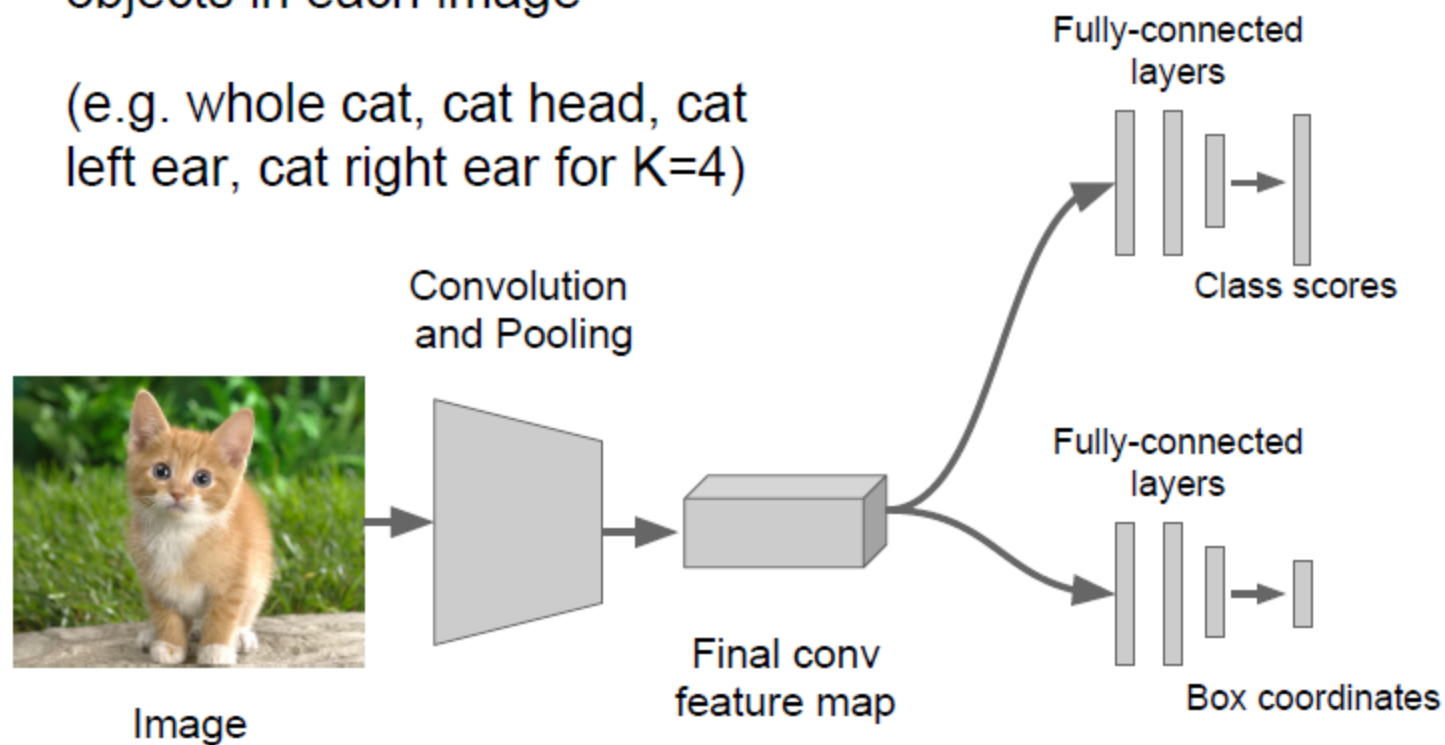
Where to attach the regression head?



Aside: Localizing multiple objects

Want to localize **exactly** K objects in each image

(e.g. whole cat, cat head, cat left ear, cat right ear for $K=4$)



$K \times 4$ numbers
(one box per object)


Aside: Human Pose Estimation

Represent a person by K joints

Regress (x, y) for each joint from last fully-connected layer of AlexNet

(Details: Normalized coordinates, iterative refinement)

Toshev and Szegedy, "DeepPose: Human Pose Estimation via Deep Neural Networks", CVPR 2014



Datasets for evaluation

- Imagenet challenges provide a platform for researchers to benchmark their novel algorithms
- PASCAL VOC 2010 is great for small scale experiments. About 1.3 GB download size.
- MS COCO datasets are available for tasks like Image Captioning. Download size is huge but selective download is possible.

| | PASCAL VOC (2010) | ImageNet Detection (ILSVRC 2014) | MS-COCO (2014) |
|--------------------------------------|-------------------------|--|-------------------|
| Number of classes | 20 | 200 | 80 |
| Number of images (train + val) | ~20k | ~470k | ~120k |
| Mean objects per image | 2.4 | 1.1 | 7.2 |